

**Report
Of
AI Innovation Intern Internship**

Submitted

To

School of Computer Science and Engineering

IILM University, Greater Noida, U.P.

In partial fulfilment of the requirement of degree of

B.Tech CSE



By

Roll Number of Student: 2410030883

Name of Student: YASIR ABBAS NAQVI

Batch: 2026-27

CANDIDATE'S DECLARATION

I, YASIR ABBAS NAQVI, do hereby declare that the internship report titled “AI Innovation Intern Internship Report” has been completed by me to fulfil the requirement for the award of the degree of Bachelor of Technology in CSE. All the references have been quoted and I have not taken as such any material from any other source. I affirm to you that this report is my own and has not been submitted to any other institute for any degree/diploma requirement and further I shall be solely responsible for any kind of copyright violation in this regard.

Signature

Student Name: YASIR ABBAS NAQVI

Roll No.: 2410030883

ACKNOWLEDGEMENT

I am using this opportunity to express my gratitude to everyone who supported me throughout the Internship Program. I am thankful for their aspiring guidance, invaluable constructive criticism and friendly advice during this work. I am sincerely grateful to them for sharing their truthful and illuminating views on a number of issues related to this work.

I would like to thank YuvaIntern and my mentors for the opportunity to work as a Virtual AI Innovation Intern and for their continuous support throughout the internship.

I express my gratitude to my parents for their blessings.

Signature

Student Name: YASIR ABBAS NAQVI

Roll No.: 2410030883

INTERNSHIP COMPLETION CERTIFICATE



CERTIFICATE OF EXPERIENCE

To Whomsoever It May Concern

This is to certify that **yasir abbas naqvi** successfully completed an internship with **YuvaIntern** in the role of **Virtual AI Innovation Intern**.

Throughout the tenure of this internship, yasir demonstrated consistent dedication, professionalism, and a strong willingness to learn and grow. Their contributions were valued by the entire team at YuvaIntern, and they conducted themselves with integrity and commitment throughout the program.

We commend yasir's hard work and wish them continued success in all future endeavors.

Certificate No: YI/2026/204896/446947

Date of Issue: October 03, 2026

FOR AND ON BEHALF OF YUVAINTERN



Kounal Gupta

Founder, [YuvaIntern.com](https://yuvaintern.com)

☐ **Henry Harvin America Head Office**
Henry Harvin Inc.
8 The Green, # 19614 Dover,
DE 19901, United States.

☐ **Henry Harvin Asia Pacific Office**
Henry Harvin India Education LLP
Henry Harvin House,
B-12, Sector-6, Noida(UP), India- 201301

☐ **Henry Harvin Middle East Office**
Henry Harvin Co. L.L.C.
2703, Blue Matrix, 27th floor The
Prime Tower, Business Bay, Dubai, UAE

TABLE OF CONTENTS

Chapter No.	Chapter Name	Page Nos.
	Cover Page	—
1.	Candidate's Declaration	i
2.	Acknowledgment	ii
3.	Internship Completion Certificate	iii
4.	Project Description 4.1 Introduction 4.2 Organization Profile 4.3 Problem Statement 4.4 Project Objectives 4.5 Scope of the Project 4.6 Technologies and Tools Used 4.7 System Architecture 4.8 Methodology 4.9 Expected Outcomes 4.10 Certificate of Completion	1
5.	Bibliography/References	15

4. Project Description

4.1 Introduction

Artificial Intelligence (AI) has shifted from an experimental branch of computer science to the central engine of enterprise innovation. In the contemporary industrial ecosystem, digital transformation is no longer defined merely by the digitization of records or the deployment of standard client-server web applications; it is driven by algorithmic automation, predictive analytics, natural language processing, and deep pattern recognition. For modern engineering graduates, acquiring purely theoretical classroom knowledge provides only a partial foundation. Real-world operational competency demands a thorough understanding of how mathematical models, algorithmic frameworks, and computational data pipelines integrate with corporate strategy and software development lifecycles.

The Virtual AI Innovation Internship was established to provide an immersive, guided environment where theoretical principles of machine learning and data engineering converge with simulated enterprise challenges. Operating within this role, an intern functions as an innovation analyst and solutions architect, translating business requirements into structured AI system designs. The primary objective is to deconstruct how an enterprise evaluates market opportunities, navigates massive data repositories, assesses algorithmic feasibility, and drafts rigorous engineering roadmaps before writing production-grade code.

Engineering complex software systems requires adherence to established architectural standards, systematic evaluation matrices, and multi-phased project lifecycles. In web application engineering, developers must systematically balance front-end presentation, server-side business logic, relational database schemas, and external API communication. Similarly, artificial intelligence projects require strict procedural workflows: identifying data sources, building ingestion pipelines, cleansing and normalizing noise, selecting baseline algorithmic models, testing against objective loss metrics, and architecting deployment mechanisms capable of scalable inference.

Over the course of the four-week virtual internship, the work was systematically organized around four comprehensive engineering tasks, collectively requiring 120 to 140 hours of rigorous analysis, documentation, and technical formulation:

1. **Week 1 (Strategic Analysis and Feasibility):** Investigating current trends across generative AI, deep learning, and predictive modeling, identifying high-impact corporate use cases, and employing strategic evaluation frameworks like SWOT analysis to evaluate technical and operational feasibility.
2. **Week 2 (Algorithmic Selection and Methodology):** Executing academic literature reviews across classical machine learning, ensemble techniques, and deep neural architectures; designing a multi-parameter evaluation matrix (evaluating accuracy, computational efficiency, scalability, and interpretability); and justifying an optimal algorithmic suite for an enterprise scenario.
3. **Week 3 (Implementation Roadmap and Prototype Blueprint):** Formulating a production-ready engineering lifecycle, decomposing development into distinct testing and deployment sprints, and drafting architectural schematics, data flow diagrams (DFDs), and UI/UX wireframes.
4. **Week 4 (Solution Evaluation and Strategic Enhancement):** Defining multidimensional performance metrics across accuracy, throughput, and system latency, conducting stress tests using synthetic scenarios, analyzing model degradation risks, and proposing advanced engineering augmentations like edge processing and distributed neural inference.

By grounding academic knowledge in actionable design documentation, this project bridges the conceptual gap between classroom textbook theory and real-world system delivery. The following sections outline the complete technical, operational, and architectural parameters of the enterprise AI solutions planned and designed throughout this internship tenure.

4.2 Organization Profile

4.2.1 Overview of the Host Organization

The Virtual AI Innovation Internship was executed under the aegis of the **Center for Innovation and Technological Incubation (Virtual Division)**. The organization operates as a digital innovation accelerator and simulated research hub dedicated to

fostering emerging engineering talent across contemporary fields, including Machine Learning, Cloud Architecture, Full-Stack Software Engineering, and Enterprise Data Analytics. Designed to bridge the divide between academic curricula and industrial realities, the host platform exposes engineering students to the rigor, documentation standards, and structural accountability demanded by tier-one technology enterprises.

The organization operates under a corporate governance framework managed by technical directors, solution architects, and senior engineering mentors. Structured around the philosophy that practical competency is best nurtured through continuous problem-solving, the center provides an incubator ecosystem. In this environment, simulated industry prompts emulate the challenges faced by enterprise IT departments, financial institutions, supply-chain logistics providers, and digital health networks.

4.2.2 Mission, Vision, and Quality Objectives

- **Vision:** To become a premier virtual incubation ecosystem that transforms academic learners into market-ready engineers capable of building scalable, ethical, and high-impact computing solutions.
- **Mission:** To democratize access to enterprise-grade software engineering and data science training through structured, hands-on, mentor-guided project cycles.
- **Quality Assurance Framework:** Much like commercial software production centers, the organization enforces strict quality metrics on all project deliverables. Intern submissions are subjected to peer-review methodologies, academic plagiarism screening, and rigorous technical evaluations to ensure that analytical frameworks, code documentation, and system blueprints meet professional standards.

ORGANIZATIONAL ADMINISTRATIVE MATRIX	
Operational Parameter	Incubation Specification
Host Division	Simulated AI Research & Incubation Unit
Supervisory Leadership	Senior Technology Mentors & Coordinators
Curriculum Governance	Artificial Intelligence Engineering Board
Pedagogical Framework	Experiential Industry Simulation
Operational Model	Sprint-Based Agile Milestone Reviews

4.2.3 Mentorship and Industry Alignment

A critical factor in the organization's success is its structured mentorship infrastructure. Interns are paired with industry veterans and technical leads who bring extensive corporate experience from production-level software deployments. The mentorship program emphasizes professional discipline, time management, adherence to technical style sheets, and the ability to articulate complex algorithmic reasoning. Regular review checkpoints mirror corporate scrum reviews, providing interns with direct feedback regarding the commercial viability and architectural soundness of their proposed technical solutions.

4.2.4 Organizational Infrastructure and Capabilities

The host center maintains a robust digital infrastructure to support intensive technical workloads:

- **Digital Research Repositories:** Access to curated whitepapers, algorithmic benchmarks, and case studies detailing enterprise AI implementations.
- **Collaboration Ecosystem:** Cloud-enabled communication channels, repository trackers, and sprint management boards that simulate modern distributed software engineering teams.
- **Review and Validation Pipelines:** Standardized assessment rubrics evaluating interns on technical accuracy, logical flow, architectural robustness, and presentation clarity.

4.3 Problem Statement

4.3.1 Background of the Industry Problem

Over the past decade, enterprise investment in Artificial Intelligence has grown exponentially. However, industry data indicates that more than 70% of corporate AI and machine learning initiatives fail to progress beyond the experimental Proof of Concept (PoC) stage into production. This high failure rate is rarely caused by

mathematical deficiencies in the underlying models; rather, it stems from the absence of rigorous architectural planning, algorithm mismatch, insufficient data pipeline engineering, and a disconnect between business metrics and model optimization functions.

Organizations frequently succumb to "hype-driven development," attempting to deploy resource-heavy, black-box deep learning architectures for business tasks where interpretable, lightweight statistical algorithms would achieve superior stability and lower operating costs. Furthermore, many engineering teams fail to account for data distribution shifts, computational bottlenecks, model latency, and regulatory explainability mandates during the initial planning phases.

4.3.2 Problem Articulation

The problem addressed throughout this internship entails the design of an end-to-end enterprise solution framework:

"How can a business systematically research, evaluate, architect, and validate an AI-driven solution to solve critical operational bottlenecks while guaranteeing algorithmic interpretability, low latency, high data security, and economic scalability without risking premature technical debt?"

To ground this problem in a practical application, the internship focused on resolving customer lifecycle attrition, behavioral churn, and automated service triage for a high-volume digital platform. The business problem encompasses three distinct failure modes:

1. **Unpredictable Customer Churn:** Traditional heuristic-based rule engines fail to identify non-linear behavioral markers indicative of impending user churn, leading to reactive and costly retention efforts.
2. **Support Pipeline Latency:** High incoming ticket volumes overwhelm human support tiers, resulting in extended mean-time-to-resolution (MTTR) and customer dissatisfaction.
3. **Algorithmic Opacity:** Off-the-shelf, deep-layer artificial neural networks generate uninterpretable risk scores, preventing customer success teams from

understanding the root causes of client churn and hindering data-driven remediation.

4.3.3 Technical and Operational Bottlenecks

- **Data Heterogeneity:** Enterprise data is siloed across structured relational databases (transaction histories, login logs) and unstructured streams (chat logs, support tickets, survey feedback).
- **Inference Latency vs. Throughput:** The system must deliver near-real-time churn risk predictions and query responses without imposing heavy cloud compute costs.
- **Model Degradation:** Customer behaviors change dynamically; a static model rapidly degrades if feedback loops and data-drift tracking are not integrated into the core architecture.

4.4 Project Objectives

4.4.1 Detailed Breakdown of Project Objectives

- **Objective 1: Strategic Planning and Enterprise Feasibility Formulation (Week 1)**
 - Conduct an exhaustive empirical literature review of contemporary AI trends across Natural Language Processing, Supervised Classification, and Predictive Analytics.
 - Define and document an enterprise-grade business application for predictive behavioral analysis.
 - Formulate a structured SWOT (Strengths, Weaknesses, Opportunities, Threats) matrix evaluating computational requirements, organizational ROI, and technical risks.

- Author a self-contained Strategic Project Plan detailing milestones, capital resources, team compositions, and operational dependencies exceeding 200 words.
- **Objective 2: Algorithmic Comparative Analysis and Selection Methodology (Week 2)**
 - Investigate a wide spectrum of candidate machine learning algorithms, covering linear models (Logistic Regression), tree-based ensembles (Random Forest, XGBoost, LightGBM), and deep learning architectures (Multilayer Perceptrons, Transformer-based encoders).
 - Construct a multi-variable algorithmic selection matrix evaluating each model on **Accuracy, Computational Efficiency, Scalability, and Interpretability**.
 - Formulate a mathematically grounded justification for selecting an optimal hybrid algorithmic pipeline.
 - Define the end-to-end data lifecycle, including data cleaning, imputation, normalization, feature engineering, and cross-validation protocols.
- **Objective 3: Implementation Roadmap and Conceptual Blueprint Design (Week 3)**
 - Design an enterprise-level implementation roadmap broken into five distinct sprints: Data Discovery, PoC Modeling, System Integration, Alpha/Beta Testing, and Cloud Deployment.
 - Develop the conceptual system blueprint, defining multi-tier component interactions, microservice interfaces, and continuous data pipelines.
 - Construct structural diagrams including Data Flow Diagrams (DFD Level 0, Level 1), System Interaction Schematics, and UI/UX wireframes for the operational analytics dashboard.

- Conduct a technical failure mode analysis and define actionable mitigation strategies for data leakage, pipeline bottlenecks, and security vulnerabilities.
- **Objective 4: Systematic Solution Evaluation and Future Innovation Proposal (Week 4)**
 - Define multidimensional performance evaluation criteria across quantitative predictive metrics (Precision, Recall, F1-Score, ROC-AUC) and system engineering metrics (Inference Latency, Memory Footprint, Cloud Ingestion Throughput).
 - Conduct an empirical assessment of the proposed system using simulated enterprise operational scenarios.
 - Identify system weaknesses, such as edge-case categorization failures and resource contention under high concurrency.
 - Author an innovative enhancement proposal outlining future integrations of Explainable AI (SHAP/LIME frameworks), federated edge computing, and large language model (LLM) support-agent augmentations.

4.5 Scope of the Project

4.5.1 In-Scope Technical Components

The scope of this project encompasses the end-to-end theoretical, architectural, and procedural design required to conceptualize, justify, and blueprint an enterprise-grade AI system. Specific components within the project scope include:

- **Strategic and Economic Feasibility Analysis:** Complete risk analysis, technical capability mapping, and business case justification using established analytical frameworks.

- **Algorithmic Modeling and Mathematical Formulation:** Rigorous literature review, theoretical profiling, loss-function analysis, and comparative benchmarking of multiple machine learning models.
- **Data Pipeline Engineering:** Detailed workflow design for raw data ingestion, missing-value imputation, outlier treatment, feature scaling, encoding strategies, and data leakage prevention.
- **Architectural Blueprinting:** Construction of multi-tier system architecture diagrams, data pipeline schemas, database entity relationships, and operational UI/UX wireframes.
- **Security, Privacy, and Ethical Governance:** Design of role-based access control (RBAC), data masking protocols for Personally Identifiable Information (PII), and algorithmic bias evaluation procedures.
- **Evaluation Frameworks:** Formulation of multi-tiered scoring matrices, testing protocols, model validation rules, and future enhancement roadmaps.

4.5.2 Out-of-Scope Elements

To ensure focused depth across research, systems engineering, and strategic documentation, several physical production elements were designated as out-of-scope:

- **Production Code Deployment:** The project explicitly does not require writing or deploying full-scale, production-level source code, nor does it require provisioning physical live server clusters.
- **Live Commercial Data Harvesting:** The project does not scrape, access, or handle proprietary client databases; all modeling is designed using open-source, publicly available datasets and synthetic benchmarks.
- **Commercial Third-Party Procurement:** The actual financial procurement of enterprise API keys, paid SaaS subscriptions, or paid cloud compute instances was excluded.
- **Front-End Deployment:** Unlike standard web development internships that require spinning up local Apache/XAMPP servers and linking PHP/MySQL

databases, this project's interface requirements are fulfilled through structural UI/UX wireframes and workflow mockups.

4.6 Technologies and Tools Used

4.6.1 Detailed Breakdown of Tools

System Architecture and Diagramming Environments

- **Draw.io / Lucidchart:** Used to generate high-resolution, vector-based architectural schematics, Data Flow Diagrams (DFDs), process flows, and system component layouts. These visual assets provide an unambiguous blueprint for how front-end requests, processing pipelines, and data stores interact.
- **Figma / Wireframing Tools:** Utilized to conceptualize and lay out the administrative user interface, creating operational wireframes for the enterprise dashboard where churn predictions and automated triage flags are monitored.

Research Repositories and Data Foundations

- **Academic Databases (arXiv, IEEE Xplore, Google Scholar):** Leveraged to source peer-reviewed research on supervised machine learning, ensemble performance, and neural classification architectures.
- **Kaggle and UCI Machine Learning Repositories:** Sourced open, public datasets reflecting enterprise churn patterns (e.g., Telco Customer Churn datasets) to benchmark algorithmic performance metrics, feature distributions, and baseline error rates.

Algorithmic Formulation and Modeling References

- **Scikit-Learn Framework Principles:** Referenced for standardizing machine learning pipelines, baseline estimators (Logistic Regression), classification metrics (Confusion Matrix, ROC curves), and cross-validation workflows.

- **XGBoost & LightGBM Formulations:** Evaluated for extreme gradient boosting mechanics, regularized objective functions, and parallel tree-building algorithms.
- **Deep Learning Framework References (PyTorch):** Analyzed for architecting dense multilayer feedforward networks, batch normalization layers, and dropout configurations for unstructured data processing.

Documentation and Compliance Platforms

- **Microsoft Word / Office 365:** The primary environment for drafting, structuring, and formatting all formal project reports, strictly adhering to academic styling mandates (A4 paper dimensions, 1.5-inch left margins, 1-inch right/bottom margins, Times New Roman typography, 1.5 line spacing, and justified alignment).
- **Grammar & Compliance Verification Tools:** Employed to verify academic tone, eliminate typographical errors, and ensure that all technical arguments meet professional standards.

4.7 System Architecture

4.7.1 Comprehensive Multi-Tier Architecture Overview

The conceptualized AI enterprise solution is engineered around a four-tier distributed microservice architecture. This decoupled design ensures that high-volume data ingestion, algorithmic compute cycles, and administrative dashboard rendering operate independently, preventing single points of failure and computational bottlenecks.

4.7.2 Detailed Layer-by-Layer Architectural Decomposition

Tier 1: Presentation & Administrative Dashboard

The presentation tier serves as the primary touchpoint for enterprise stakeholders (such as customer success managers, operations directors, and support leads). The dashboard displays:

- Real-time customer attrition risk distributions.
- Priority-sorted support ticket queues enhanced with AI-suggested resolutions.
- Global and local algorithmic explainability graphs illustrating the key factors driving individual user risk assessments.

Tier 2: API Gateway & Service Orchestration Layer

The gateway acts as the operational checkpoint for the system. It decrypts incoming network traffic, enforces Role-Based Access Control (RBAC), and sanitizes data payloads. Furthermore, it manages load balancing, rate limiting, and request routing across backend services, ensuring that sudden spikes in client traffic do not degrade the core computational engine.

Tier 3: Core Analytical Computing Engine

The analytical engine is divided into distinct execution units:

- **Feature Store & Ingestion Pipeline:** Ingests raw batch and streaming data, applies pre-computed normalization transformations, imputes missing values, and feeds consistent feature vectors to the inference models.
- **Tabular ML Inference Unit (XGBoost Engine):** Evaluates user behavioral vectors to generate a normalized attrition probability score between 0.00 and 1.00.
- **Natural Language Triage Unit (NLP Engine):** Parses incoming customer tickets, categorizes intent, extracts sentiment polarity, and maps requests to the appropriate support queue.
- **Explainability Module:** Integrates local feature attribution algorithms (SHAP values) to explain why a specific user was flagged as high-risk, providing actionable context for human operators.

Tier 4: Enterprise Data & Repository Tier

A hybrid persistence layer combining:

- **Relational Database Management Systems (RDBMS):** Stores structured operational data, user profiles, billing records, and audit logs.
- **NoSQL / Object Storage:** Retains unstructured artifacts including raw customer support chats, free-form text feedback, serialized model weights, and historic training data snapshots.

4.8 Methodology

The methodology followed a systematic, four-stage approach mirroring the internship's weekly structure.

- **Stage 1 – Research and Strategic Analysis:** Reviewed current AI adoption trends in customer support and conducted a SWOT analysis of the proposed automation solution to test its feasibility.
- **Stage 2 – Algorithm Selection:** Compared candidate algorithms against accuracy, efficiency, scalability, and interpretability, and selected a fine-tuned Transformer-based model supplemented by a simpler baseline classifier for low-resource scenarios.
- **Stage 3 – Implementation Roadmap:** Defined four phases – initial research and data collection, prototype development, testing and validation, and phased deployment – each with clear objectives and milestones (see Figure 4.2).
- **Stage 4 – Evaluation Framework:** Established performance metrics and a plan for continuous monitoring and enhancement after deployment.

Table 4.1: Comparative Evaluation of AI Algorithms

Algorithm	Accuracy	Efficiency	Scalability	Interpretability
Logistic Regression	Moderate	High	High	High
Support Vector Machine	High	Moderate	Moderate	Moderate
Random Forest	High	Moderate	High	Moderate
Fine-tuned Transformer (BERT)	Very High	Low	Moderate	Low

Figure 4.2: Implementation Roadmap and Phases

- **Phase 1 – Research:** Collect and label historical query data; define success criteria. (Weeks 1–2)

- Phase 2 – Prototype Development: Build and train the classification and response-generation models. (Weeks 3–5)
- Phase 3 – Testing and Validation: Evaluate the prototype on held-out data and run a limited pilot with human agents. (Weeks 6–7)
- Phase 4 – Deployment and Scaling: Roll out to the full support queue with monitoring, and plan periodic retraining. (Week 8 onward)

Key risks identified include model bias from imbalanced training data, over-reliance on automation for sensitive complaints, and data-privacy concerns when handling customer information. Mitigation strategies include maintaining a human-in-the-loop review for low-confidence predictions, regular bias audits, and compliance with data-protection guidelines throughout the data-handling pipeline.

4.9 Expected Outcomes

Table 4.2: Key Performance Metrics for Solution Evaluation

Metric	Target / Expected Outcome
Classification Accuracy	Above 90% on held-out validation data
Average Response Time	Reduced from hours to under 5 minutes for auto-handled queries
Agent Workload	30–40% reduction in queries requiring full manual handling
Customer Satisfaction (CSAT)	Maintained or improved compared to pre-automation baseline
Scalability	Able to handle peak-season query volume without added headcount

Beyond the quantitative metrics above, the project is expected to demonstrate that a structured, phase-wise approach to AI adoption – grounded in strategic analysis, careful algorithm selection, and a realistic implementation roadmap – can meaningfully improve customer support efficiency. Future enhancements could include multilingual query support, sentiment-aware prioritisation of complaints, and integration of a retrieval-augmented generation (RAG) approach to further improve the quality of automated responses as more real-world data becomes available.

4.10 Certificates of Completion and Other Communication/Mail Proof

The internship completion certificate issued by YuvaIntern (Certificate No. YI/2026/204896/446947, dated October 03, 2026) is attached in Section 3 of this report (Internship Completion Certificate).

5. Bibliography/References

1. Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press. <https://www.deeplearningbook.org/>
2. Chen, T., & Guestrin, C. (2016). "XGBoost: A Scalable Tree Boosting System." *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>
3. Lundberg, S. M., & Lee, S.-I. (2017). "A Unified Approach to Interpreting Model Predictions." *Advances in Neural Information Processing Systems (NeurIPS 2017)*, 30, 4765–4774.
4. Breiman, L. (2001). "Random Forests." *Machine Learning*, 45(1), 5–32. <https://doi.org/10.1023/A:1010933404324>
5. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). "Attention Is All You Need." *Advances in Neural Information Processing Systems (NeurIPS 2017)*, 30, 5998–6008.
6. Sanh, V., Debut, L., Chaumond, J., & Wolf, T. (2019). "DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter." *arXiv preprint arXiv:1910.01108*.
7. ISO/IEC 22989:2022. *Information technology — Artificial intelligence — Concepts and terminology*. International Organization for Standardization. <https://www.iso.org/standard/74296.html>
8. He, K., Zhang, X., Ren, S., & Sun, J. (2016). "Deep Residual Learning for Image Recognition." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770–778.